

Class 13: RNASeq Analysis with DESq2

PID:A16631132 Nashed

The data for this hands-on session comes from a published RNA-seq experiment where airway smooth muscle cells were treated with **dexamethasone** (dex), a synthetic glucocorticoid steroid with anti-inflammatory effects (Himes et al. 2014).

```
# Complete the missing code
counts <- read.csv("airway_scaledcounts.csv", row.names=1)
metadata <- read.csv("airway_metadata.csv")
```

```
head(counts)
```

| | SRR1039508 | SRR1039509 | SRR1039512 | SRR1039513 | SRR1039516 |
|------------------|------------|------------|------------|------------|------------|
| ENSG000000000003 | 723 | 486 | 904 | 445 | 1170 |
| ENSG000000000005 | 0 | 0 | 0 | 0 | 0 |
| ENSG000000000419 | 467 | 523 | 616 | 371 | 582 |
| ENSG000000000457 | 347 | 258 | 364 | 237 | 318 |
| ENSG000000000460 | 96 | 81 | 73 | 66 | 118 |
| ENSG000000000938 | 0 | 0 | 1 | 0 | 2 |

| | SRR1039517 | SRR1039520 | SRR1039521 |
|------------------|------------|------------|------------|
| ENSG000000000003 | 1097 | 806 | 604 |
| ENSG000000000005 | 0 | 0 | 0 |
| ENSG000000000419 | 781 | 417 | 509 |
| ENSG000000000457 | 447 | 330 | 324 |
| ENSG000000000460 | 94 | 102 | 74 |
| ENSG000000000938 | 0 | 0 | 0 |

Q1. How many genes are in this dataset?

```
nrow(counts)
```

```
[1] 38694
```

Q2. How many ‘control’ cell lines do we have?

```
table(metadata$dex)
```

```
control treated  
      4      4
```

```
sum(metadata$dex=="control")
```

```
[1] 4
```

Toy differential gene expression

Let’s start by calculating the mean counts per gene in the “control” samples. We can then compare this value for each gene to the mean counts in the “treated” samples (i.e. columns).

- Step 1. Find which columns in the counts correspond to “control” samples.
- Step 2. Calculate the mean value per gene in these columns.
- Step 3. Store my answer for later in `control.mean`

```
head(counts)
```

| | SRR1039508 | SRR1039509 | SRR1039512 | SRR1039513 | SRR1039516 |
|------------------|------------|------------|------------|------------|------------|
| ENSG000000000003 | 723 | 486 | 904 | 445 | 1170 |
| ENSG000000000005 | 0 | 0 | 0 | 0 | 0 |
| ENSG000000000419 | 467 | 523 | 616 | 371 | 582 |
| ENSG000000000457 | 347 | 258 | 364 | 237 | 318 |
| ENSG000000000460 | 96 | 81 | 73 | 66 | 118 |
| ENSG000000000938 | 0 | 0 | 1 | 0 | 2 |

| | SRR1039517 | SRR1039520 | SRR1039521 |
|------------------|------------|------------|------------|
| ENSG000000000003 | 1097 | 806 | 604 |
| ENSG000000000005 | 0 | 0 | 0 |
| ENSG000000000419 | 781 | 417 | 509 |
| ENSG000000000457 | 447 | 330 | 324 |
| ENSG000000000460 | 94 | 102 | 74 |
| ENSG000000000938 | 0 | 0 | 0 |

```
control.inds <- metadata$dex == "control"
```

```
metadata[control.inds,]
```

```

      id      dex celltype      geo_id
1 SRR1039508 control   N61311 GSM1275862
3 SRR1039512 control   N052611 GSM1275866
5 SRR1039516 control   N080611 GSM1275870
7 SRR1039520 control   N061011 GSM1275874

```

```

control.counts <- counts[,control.inds]
head(control.counts)

```

| | SRR1039508 | SRR1039512 | SRR1039516 | SRR1039520 |
|-------------------|------------|------------|------------|------------|
| ENSG000000000003 | 723 | 904 | 1170 | 806 |
| ENSG000000000005 | 0 | 0 | 0 | 0 |
| ENSG0000000000419 | 467 | 616 | 582 | 417 |
| ENSG0000000000457 | 347 | 364 | 318 | 330 |
| ENSG0000000000460 | 96 | 73 | 118 | 102 |
| ENSG0000000000938 | 0 | 1 | 2 | 0 |

Q3. How would you make the above code in either approach more robust? Is there a function that could help here? You can use a code that does not state the number of samples, like `RowMeans`

```

#apply(control.counts, 1, mean)
control.mean <- rowMeans(control.counts)

```

Q4. Follow the same procedure for the treated samples (i.e. calculate the mean per gene across drug treated samples and assign to a labeled vector called `treated.mean`)

```
treated.mean <- rowMeans (counts[,metadata$dex == "treated"])
```

To keep us tidy lets put `control.mean` and `treated.mean` vectors together as two columns of a new data.frame. `meancounts <- data.frame(control.mean, treated.mean)`


```
meancounts <- data.frame(control.mean, treated.mean)
```

```
head(meancounts)
```

| | control.mean | treated.mean |
|------------------|--------------|--------------|
| ENSG000000000003 | 900.75 | 658.00 |
| ENSG000000000005 | 0.00 | 0.00 |
| ENSG000000000419 | 520.50 | 546.00 |
| ENSG000000000457 | 339.75 | 316.50 |
| ENSG000000000460 | 97.25 | 78.75 |
| ENSG000000000938 | 0.75 | 0.00 |

Q5 (a). Create a scatter plot showing the mean of the treated samples against the mean of the control samples. Your plot should look something like the following.


```
plot(meancounts)
```



Q5 (b). You could also use the ggplot2 package to make this figure producing the plot below. What geom_?() function would you use for this plot? a ggplot version:

```
library(ggplot2)

ggplot(meancounts) +
  aes(control.mean, treated.mean) +
  geom_point(alpha=0.2)
```




Q6. Try plotting both axes on a log scale. What is the argument to `plot()` that allows you to do this?

```
plot(meancounts, log="xy")
```

```
Warning in xy.coords(x, y, xlabel, ylabel, log): 15032 x values <= 0 omitted  
from logarithmic plot
```

```
Warning in xy.coords(x, y, xlabel, ylabel, log): 15281 y values <= 0 omitted  
from logarithmic plot
```



Log transformation are super useful when our data is skewed and measured over a wide range like this. We can use different log transformations like base10 or natural logs butt we most often prefer log2 units.

```
#Treated/Control
log2(10/10)
```

```
[1] 0
```

What if there was a doubling

```
# Treated/Control
log2(20/10)
```

```
[1] 1
```

Half counts

```
log2( 10/20)
```

```
[1] -1
```

```
log2(40/10)
```

```
[1] 2
```

```
log10(40/10)
```

```
[1] 0.60206
```

Let's add a log2 fold-change column to our little `mean.counts` data.frame:

```
meancounts$log2fc <- log2(meancounts$treated.mean/  
                           meancounts$control.mean)  
head(meancounts)
```

| | control.mean | treated.mean | log2fc |
|------------------|--------------|--------------|-------------|
| ENSG000000000003 | 900.75 | 658.00 | -0.45303916 |
| ENSG000000000005 | 0.00 | 0.00 | NaN |
| ENSG000000000419 | 520.50 | 546.00 | 0.06900279 |
| ENSG000000000457 | 339.75 | 316.50 | -0.10226805 |
| ENSG000000000460 | 97.25 | 78.75 | -0.30441833 |
| ENSG000000000938 | 0.75 | 0.00 | -Inf |

There are a couple of weird results. Namely, the NaN (not a number) and -Inf results. The NaN is returned when you divide by zero and try to take the log. The -Inf is returned when you try to take the log of zero. It turns out that there are a lot of genes with zero expression. Let's filter our data to remove these genes. Again inspect your result (and the intermediate steps) to see if things make sense to you

```
to.rm.inds <- rowSums( meancounts[,1:2] == 0 ) > 0  
mycounts <- meancounts [ !to.rm.inds, ]
```

The `!` mark flips TRUE values to FALSE and vice-versa...

```
x <- c(TRUE, FALSE, TRUE)  
!x
```

```
[1] FALSE TRUE FALSE
```

```
x
```

```
[1] TRUE FALSE TRUE
```

```
which(x)
```

```
[1] 1 3
```

```
dim(mycounts)
```

```
[1] 21817      3
```

```
head(mycounts)
```

| | control.mean | treated.mean | log2fc |
|------------------|--------------|--------------|-------------|
| ENSG000000000003 | 900.75 | 658.00 | -0.45303916 |
| ENSG000000000419 | 520.50 | 546.00 | 0.06900279 |
| ENSG000000000457 | 339.75 | 316.50 | -0.10226805 |
| ENSG000000000460 | 97.25 | 78.75 | -0.30441833 |
| ENSG000000000971 | 5219.00 | 6687.50 | 0.35769358 |
| ENSG000000001036 | 2327.00 | 1785.75 | -0.38194109 |

A common threshold used for calling something differentially expressed is a $\log_2(\text{FoldChange})$ of greater than 2 or less than -2.

Let's filter the dataset both ways to see how many genes are up or down-regulated.

Q7. . What is the purpose of the `arr.ind` argument in the `which()` function call above? Why would we then take the first column of the output and need to call the `unique()` function? The `arr.ind` argument helps the `which` output the row and column positions and if you use it with `TRUE`, it will only give the true values and give the ones that have 0. If you use `unique()` it helps to not count the same row two times.

```
up.ind <- mycounts$log2fc > 2  
down.ind <- mycounts$log2fc < (-2)
```

Q8. Using the `up.ind` vector above can you determine how many up regulated genes we have at the greater than 2 fc level?


```
sum(up.ind)
```

```
[1] 250
```

Q9. Using the down.ind vector above can you determine how many down regulated genes we have at the greater than 2 fc level?

```
sum(down.ind)
```

```
[1] 367
```

Q10. Do you trust these results? Why or why not? There could be huge variance and we do not have significance, we need to know if that difference could be significant.

But we forgot all about statistical significance of these differences...

We will use the DESeq2 package to do this analysis properly...

Using DESeq2

Like any package we must load it up with a `library()` call.

```
library(DESeq2)
```

```
Loading required package: S4Vectors
```

```
Loading required package: stats4
```

```
Loading required package: BiocGenerics
```

```
Attaching package: 'BiocGenerics'
```

```
The following objects are masked from 'package:stats':
```

```
IQR, mad, sd, var, xtabs
```

The following objects are masked from 'package:base':

```
anyDuplicated, aperm, append, as.data.frame, basename, cbind,  
colnames, dirname, do.call, duplicated, eval, evalq, Filter, Find,  
get, grep, grepl, intersect, is.unsorted, lapply, Map, mapply,  
match, mget, order, paste, pmax, pmax.int, pmin, pmin.int,  
Position, rank, rbind, Reduce, rownames, sapply, setdiff, sort,  
table, tapply, union, unique, unsplit, which.max, which.min
```

Attaching package: 'S4Vectors'

The following object is masked from 'package:utils':

```
findMatches
```

The following objects are masked from 'package:base':

```
expand.grid, I, unname
```

Loading required package: IRanges

Loading required package: GenomicRanges

Loading required package: GenomeInfoDb

Loading required package: SummarizedExperiment

Loading required package: MatrixGenerics

Loading required package: matrixStats

Attaching package: 'MatrixGenerics'

The following objects are masked from 'package:matrixStats':

```
colAlls, colAnyNAs, colAnys, colAvgPerRowSet, colCollapse,
colCounts, colCummaxs, colCummins, colCumprods, colCumsums,
colDiffs, colIQRDiffs, colIQRs, colLogSumExps, colMadDiffs,
colMads, colMaxs, colMeans2, colMedians, colMins, colOrderStats,
colProds, colQuantiles, colRanges, colRanks, colSdDiffs, colSds,
colSums2, colTabulates, colVarDiffs, colVars, colWeightedMads,
colWeightedMeans, colWeightedMedians, colWeightedSds,
colWeightedVars, rowAlls, rowAnyNAs, rowAnys, rowAvgPerColSet,
rowCollapse, rowCounts, rowCummaxs, rowCummins, rowCumprods,
rowCumsums, rowDiffs, rowIQRDiffs, rowIQRs, rowLogSumExps,
rowMadDiffs, rowMads, rowMaxs, rowMeans2, rowMedians, rowMins,
rowOrderStats, rowProds, rowQuantiles, rowRanges, rowRanks,
rowSdDiffs, rowSds, rowSums2, rowTabulates, rowVarDiffs, rowVars,
rowWeightedMads, rowWeightedMeans, rowWeightedMedians,
rowWeightedSds, rowWeightedVars
```

Loading required package: Biobase

Welcome to Bioconductor

```
Vignettes contain introductory material; view with
'browseVignettes()'. To cite Bioconductor, see
'citation("Biobase")', and for packages 'citation("pkgname")'.
```

Attaching package: 'Biobase'

The following object is masked from 'package:MatrixGenerics':

```
rowMedians
```

The following objects are masked from 'package:matrixStats':

```
anyMissing, rowMedians
```

Setup the input object required by DESeq

```
dds <- DESeqDataSetFromMatrix(countData = counts,
                              colData= metadata,
                              design = ~dex)
```

converting counts to integer mode

Warning in DESeqDataSet(se, design = design, ignoreRank): some variables in design formula are characters, converting to factors

Now we can run our DESeq analysis

```
dds <- DESeq(dds)
```

estimating size factors

estimating dispersions

gene-wise dispersion estimates

mean-dispersion relationship

final dispersion estimates

fitting model and testing

Get our results back from the dds object.

```
res <- results(dds)
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 6 columns

| | baseMean | log2FoldChange | lfcSE | stat | pvalue |
|------------------|------------|----------------|-----------|-----------|-----------|
| | <numeric> | <numeric> | <numeric> | <numeric> | <numeric> |
| ENSG000000000003 | 747.194195 | -0.3507030 | 0.168246 | -2.084470 | 0.0371175 |
| ENSG000000000005 | 0.000000 | NA | NA | NA | NA |


| | | | | | |
|------------------|------------|------------|----------|-----------|-----------|
| ENSG000000000419 | 520.134160 | 0.2061078 | 0.101059 | 2.039475 | 0.0414026 |
| ENSG000000000457 | 322.664844 | 0.0245269 | 0.145145 | 0.168982 | 0.8658106 |
| ENSG000000000460 | 87.682625 | -0.1471420 | 0.257007 | -0.572521 | 0.5669691 |
| ENSG000000000938 | 0.319167 | -1.7322890 | 3.493601 | -0.495846 | 0.6200029 |
| | padj | | | | |
| | <numeric> | | | | |
| ENSG000000000003 | 0.163035 | | | | |
| ENSG000000000005 | NA | | | | |
| ENSG000000000419 | 0.176032 | | | | |
| ENSG000000000457 | 0.961694 | | | | |
| ENSG000000000460 | 0.815849 | | | | |
| ENSG000000000938 | NA | | | | |

#A Summary results plot

Volcano plot.

This is a common type of summary figure that keeps both our inner biologist and inner stats nerd happy because it shows both P-values and log2 (fold-changes)

```
plot(res$log2FoldChange, -log(res$padj))
abline(v=2, col="red")
abline(v=-2, col="red")
abline(h=-log(0.05), col="blue")
```



```
log(0.1)
```

```
[1] -2.302585
```

```
log(0.00001)
```

```
[1] -11.51293
```


```
# Setup our custom point color vector
mycols <- rep("gray", nrow(res))
mycols[ abs(res$log2FoldChange) > 2 ] <- "red"

inds <- (res$padj < 0.01) & (abs(res$log2FoldChange) > 2 )
mycols[ inds ] <- "blue"

# Volcano plot with custom colors
plot( res$log2FoldChange, -log(res$padj),
      col=mycols, ylab="-Log(P-value)", xlab="Log2(FoldChange)" )

# Cut-off lines
```

```
abline(v=c(-2,2), col="gray", lty=2)
abline(h=-log(0.1), col="gray", lty=2)
```



Save our results to date.....

```
write.csv(res, file= "deseq_results.csv")
```

```
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 6 columns

| | baseMean | log2FoldChange | lfcSE | stat | pvalue |
|------------------|------------|----------------|-----------|-----------|-----------|
| | <numeric> | <numeric> | <numeric> | <numeric> | <numeric> |
| ENSG000000000003 | 747.194195 | -0.3507030 | 0.168246 | -2.084470 | 0.0371175 |
| ENSG000000000005 | 0.000000 | NA | NA | NA | NA |
| ENSG000000000419 | 520.134160 | 0.2061078 | 0.101059 | 2.039475 | 0.0414026 |
| ENSG000000000457 | 322.664844 | 0.0245269 | 0.145145 | 0.168982 | 0.8658106 |
| ENSG000000000460 | 87.682625 | -0.1471420 | 0.257007 | -0.572521 | 0.5669691 |
| ENSG000000000938 | 0.319167 | -1.7322890 | 3.493601 | -0.495846 | 0.6200029 |

padj

| | <numeric> |
|------------------|-----------|
| ENSG000000000003 | 0.163035 |
| ENSG000000000005 | NA |
| ENSG000000000419 | 0.176032 |
| ENSG000000000457 | 0.961694 |
| ENSG000000000460 | 0.815849 |
| ENSG000000000938 | NA |

Adding Annotation data

```
library("AnnotationDbi")
```

Warning: package 'AnnotationDbi' was built under R version 4.3.2

```
library("org.Hs.eg.db")
```

```
org.Hs.eg.db
```

OrgDb object:

```
| DBSCHEMAVERSION: 2.1
| Db type: OrgDb
| Supporting package: AnnotationDbi
| DBSCHEMA: HUMAN_DB
| ORGANISM: Homo sapiens
| SPECIES: Human
| EGSOURCEDATE: 2023-Sep11
| EGSOURCENAME: Entrez Gene
| EGSOURCEURL: ftp://ftp.ncbi.nlm.nih.gov/gene/DATA
| CENTRALID: EG
| TAXID: 9606
| GOSOURCENAME: Gene Ontology
| GOSOURCEURL: http://current.geneontology.org/ontology/go-basic.obo
| GOSOURCEDATE: 2023-07-27
| GOEGSOURCEDATE: 2023-Sep11
| GOEGSOURCENAME: Entrez Gene
| GOEGSOURCEURL: ftp://ftp.ncbi.nlm.nih.gov/gene/DATA
```



```
| KEGGSOURCENAME: KEGG GENOME
| KEGGSOURCEURL: ftp://ftp.genome.jp/pub/kegg/genomes
| KEGGSOURCEDATE: 2011-Mar15
| GPSOURCENAME: UCSC Genome Bioinformatics (Homo sapiens)
| GPSOURCEURL:
| GPSOURCEDATE: 2023-Aug20
| ENSOURCEDATE: 2023-May10
| ENSOURCENAME: Ensembl
| ENSOURCEURL: ftp://ftp.ensembl.org/pub/current_fasta
| UPSOURCENAME: Uniprot
| UPSOURCEURL: http://www.UniProt.org/
| UPSOURCEDATE: Mon Sep 18 16:12:39 2023
```

Please see: `help('select')` for usage information

```
columns(org.Hs.eg.db)
```

```
[1] "ACCNUM"      "ALIAS"      "ENSEMBL"    "ENSEMBLPROT" "ENSEMBLTRANS"
[6] "ENTREZID"    "ENZYME"     "EVIDENCE"   "EVIDENCEALL" "GENENAME"
[11] "GENETYPE"    "GO"         "GOALL"      "IPI"         "MAP"
[16] "OMIM"        "ONTOLOGY"   "ONTOLOGYALL" "PATH"        "PFAM"
[21] "PMID"        "PROSITE"    "REFSEQ"     "SYMBOL"      "UCSCKG"
[26] "UNIPROT"
```

Our current IDs are here: The main function we will use here is called `mapIds()`

```
#mapIds()
head(row.names(res))
```

```
[1] "ENSG000000000003" "ENSG000000000005" "ENSG000000000419" "ENSG000000000457"
[5] "ENSG000000000460" "ENSG000000000938"
```

These are in ENSEMBLE format. I want “SYMBOL” ids.

```
res$symbol <- mapIds(org.Hs.eg.db,
                     keys=row.names(res), # Our genenames
                     keytype="ENSEMBL",   # The format of our genenames
                     column="SYMBOL",      # The new format we want to add
```

```
multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 7 columns

| | baseMean | log2FoldChange | lfcSE | stat | pvalue |
|-------------------|------------|----------------|-----------|-----------|-----------|
| | <numeric> | <numeric> | <numeric> | <numeric> | <numeric> |
| ENSG000000000003 | 747.194195 | -0.3507030 | 0.168246 | -2.084470 | 0.0371175 |
| ENSG000000000005 | 0.000000 | NA | NA | NA | NA |
| ENSG0000000000419 | 520.134160 | 0.2061078 | 0.101059 | 2.039475 | 0.0414026 |
| ENSG0000000000457 | 322.664844 | 0.0245269 | 0.145145 | 0.168982 | 0.8658106 |
| ENSG0000000000460 | 87.682625 | -0.1471420 | 0.257007 | -0.572521 | 0.5669691 |
| ENSG0000000000938 | 0.319167 | -1.7322890 | 3.493601 | -0.495846 | 0.6200029 |
| | padj | symbol | | | |
| | <numeric> | <character> | | | |
| ENSG0000000000003 | 0.163035 | TSPAN6 | | | |
| ENSG0000000000005 | NA | TNMD | | | |
| ENSG0000000000419 | 0.176032 | DPM1 | | | |
| ENSG0000000000457 | 0.961694 | SCYL3 | | | |
| ENSG0000000000460 | 0.815849 | FIRRM | | | |
| ENSG0000000000938 | NA | FGR | | | |

```
res$entrez <- mapIds(org.Hs.eg.db,
                      keys=row.names(res), # Our genenames
                      keytype="ENSEMBL",   # The format of our genenames
                      column="ENTREZID",   # The new format we want to add
                      multiVals="first")
```

'select()' returned 1:many mapping between keys and columns

```
head(res)
```

log2 fold change (MLE): dex treated vs control

Wald test p-value: dex treated vs control

DataFrame with 6 rows and 8 columns

| | baseMean | log2FoldChange | lfcSE | stat | pvalue |
|------------------|------------|----------------|-------------|-----------|-----------|
| | <numeric> | <numeric> | <numeric> | <numeric> | <numeric> |
| ENSG000000000003 | 747.194195 | -0.3507030 | 0.168246 | -2.084470 | 0.0371175 |
| ENSG000000000005 | 0.000000 | NA | NA | NA | NA |
| ENSG000000000419 | 520.134160 | 0.2061078 | 0.101059 | 2.039475 | 0.0414026 |
| ENSG000000000457 | 322.664844 | 0.0245269 | 0.145145 | 0.168982 | 0.8658106 |
| ENSG000000000460 | 87.682625 | -0.1471420 | 0.257007 | -0.572521 | 0.5669691 |
| ENSG000000000938 | 0.319167 | -1.7322890 | 3.493601 | -0.495846 | 0.6200029 |
| | padj | symbol | entrez | | |
| | <numeric> | <character> | <character> | | |
| ENSG000000000003 | 0.163035 | TSPAN6 | 7105 | | |
| ENSG000000000005 | NA | TNMD | 64102 | | |
| ENSG000000000419 | 0.176032 | DPM1 | 8813 | | |
| ENSG000000000457 | 0.961694 | SCYL3 | 57147 | | |
| ENSG000000000460 | 0.815849 | FIRRM | 55732 | | |
| ENSG000000000938 | NA | FGR | 2268 | | |

##Pathway Analysis

We will use the **gage** package along with **pathview** here to do genset enrichment (a.k.a pathway analysis) and figure generation respectively.

```
#1 message: false
library(pathview)
```

```
#####
Pathview is an open source software package distributed under GNU General
Public License version 3 (GPLv3). Details of GPLv3 is available at
http://www.gnu.org/licenses/gpl-3.0.html. Particullary, users are required to
formally cite the original Pathview paper (not just mention it) in publications
or products. For details, do citation("pathview") within R.
```

The pathview downloads and uses KEGG data. Non-academic uses may require a KEGG license agreement (details at <http://www.kegg.jp/kegg/legal.html>).

```
#####
```

```
library(gage)
```

```
library(gageData)
```

Lets have a peak at the first two pathways in KEGG

```
data(kegg.sets.hs)
```

```
# Examine the first 2 pathways in this kegg set for humans
head(kegg.sets.hs, 2)
```

```
$`hsa00232 Caffeine metabolism`
```

```
[1] "10" "1544" "1548" "1549" "1553" "7498" "9"
```

```
$`hsa00983 Drug metabolism - other enzymes`
```

```
[1] "10" "1066" "10720" "10941" "151531" "1548" "1549" "1551"
[9] "1553" "1576" "1577" "1806" "1807" "1890" "221223" "2990"
[17] "3251" "3614" "3615" "3704" "51733" "54490" "54575" "54576"
[25] "54577" "54578" "54579" "54600" "54657" "54658" "54659" "54963"
[33] "574537" "64816" "7083" "7084" "7172" "7363" "7364" "7365"
[41] "7366" "7367" "7371" "7372" "7378" "7498" "79799" "83549"
[49] "8824" "8833" "9" "978"
```

What we need for `gage()` is our genes in ENTREZ id format with a measure of their importance.

It wants a vector of e.g. fold-changes.

```
foldchanges <- res$log2FoldChange
head(foldchanges)
```

```
[1] -0.35070302 NA 0.20610777 0.02452695 -0.14714205 -1.73228897
```

```
x <- c(100, 80, 100)
names(x) <- c("desteny", "barry", "chris")
x
```

```
desteny barry chris
100      80      100
```

Add ENTREZ ids as `names()` to my `foldchanges` vector.

```
names(foldchanges) <- res$entrez
head(foldchanges)
```

```
      7105      64102      8813      57147      55732      2268
-0.35070302      NA  0.20610777  0.02452695 -0.14714205 -1.73228897
```

Now we can run `gage()` with this input vector and the gneset we want to examine for overlap/enrichment...

```
# Get the results
keggres = gage(foldchanges, gsets=kegg.sets.hs)
```

Look at the results.

```
attributes(keggres)
```

```
$names
[1] "greater" "less"    "stats"
```

```
head(keggres$less, 3)
```

| | | p.geomean | stat.mean | p.val |
|----------|---------------------------|--------------|-----------|--------------|
| hsa05332 | Graft-versus-host disease | 0.0004250461 | -3.473346 | 0.0004250461 |
| hsa04940 | Type I diabetes mellitus | 0.0017820293 | -3.002352 | 0.0017820293 |
| hsa05310 | Asthma | 0.0020045888 | -3.009050 | 0.0020045888 |

| | | q.val | set.size | exp1 |
|----------|---------------------------|------------|----------|--------------|
| hsa05332 | Graft-versus-host disease | 0.09053483 | 40 | 0.0004250461 |
| hsa04940 | Type I diabetes mellitus | 0.14232581 | 42 | 0.0017820293 |
| hsa05310 | Asthma | 0.14232581 | 29 | 0.0020045888 |

We can view these pathways with our geneset genes highlighted using the `pathview()` function. E.g. for “Asthma” I will use the pathway.id hsa05310 as seen above.

```
pathview(gene.data=foldchanges, pathway.id="hsa05310")
```

'select()' returned 1:1 mapping between keys and columns

Info: Working in directory /Users/nicolenashed/Desktop/R Coding/class13

Info: Writing image file hsa05310.pathview.png

